

Deep Learning

Lecture 8

Data Augmentation

Pretrained Models and

Transfer Learning

Prof. Dr. Rainer Herrler

Phone: 09721/ 940-8710

Email: rainer.herrler@fhws.de

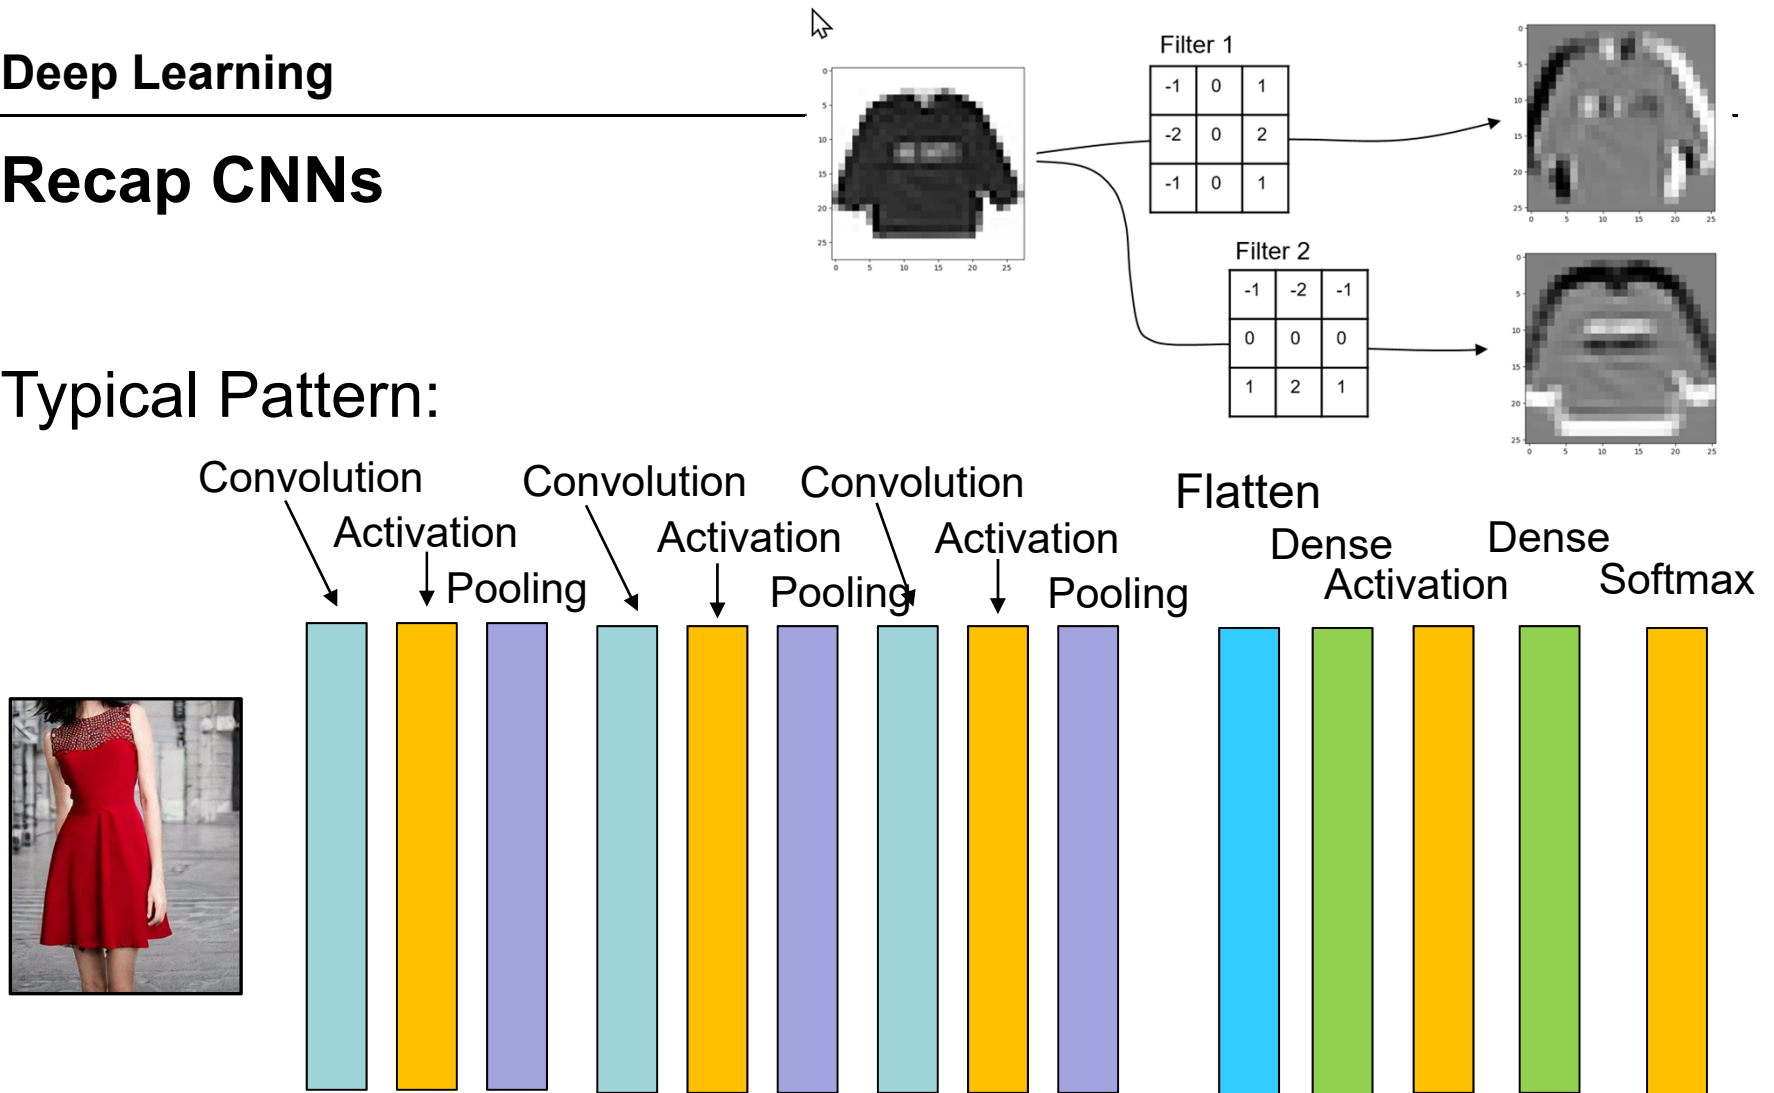
Pictures from Wikipedia / Pixabay

Some Pictures generated with Stable Diffusion

Deep Learning

Recap CNNs


Typical Pattern:



Other layers like Normalization / Dropout etc can also be inserted

Question: How can we see what's generated from the first layer?

GPU's for Deep Learning



GPU Architectures /Generations:

- Kepler, Pascal, Volta (2017), Turing (2018), Ampere (2020), Ada & Hopper (2022)

Use Case:

- Tegra (Embedded) / GeForce (Consumer)
- Quadro (Workstation) / Tesla (Datacenter)

e.g. my Notebook GTX 1650 mobile is a Turing Architecture ~6x slower than T4

Overview on current hosted services including GPU

Free hosted Jupiter Notebooks

- Google Colab
 - <https://colab.research.google.com/>
 - Direct connection to google drive for persisting data (model weights, training data)
 - T4 GPU or comparable TPU
 - Upgradeable to Colab Pro for stonger GPUs
- Kaggle
 - <https://www.kaggle.com/>
 - Good for exercises with data provided by Kaggle
 - Provides P100 and T4 GPUs as well as TPU's
- Paperspace
 - <https://www.paperspace.com/gpu-cloud>
 - Free GPUs of different types, even very powerful ones
 - Very often out of resources

Payed Services

- Microsoft Azure
 - Free 100\$ budget for students
 - Many Services provided –e.g. Azure Machine Learning Studio
 - Complex setup
 - Eg 90cent per h for Tesla K80
- Lambda Labs
 - <https://lambdalabs.com/service/gpu-cloud/pricing>
 - Newest types of GPU,
 - Good prices starting from 0.5\$ per h
- Vast.ai
 - <https://cloud.vast.ai/create/>
 - Very Cheap options available,
 - From private or Datacenter
- Amazon, Jarvislabs...

Exercise



Try to create a CNN for CIFAR 100

- Short Link: <https://bit.ly/3Oj1YeG>
- Full Link: https://colab.research.google.com/drive/1x_PCB1JdyLZgG17nJ8ePBoO95RO22C7C?usp=sharing
- Difficulties
 - Many classes and few samples
 - Complex networks will overfit very fast
 - Regularization is needed!
- Realistic outcome:
 - Rates around >35 % accuracy should be achievable
 - More complex networks can reach 60% accuracy and more

[DL_015_CIFAR-100-CNNs.ipynb](#)



Further comments when training in the cloud

- Save the training results regularly
 - Typically your notebooks are persistently stored
 - Training data and trained weights are **not on a (very) permanent store**
 - Store training data somewhere else (Github/Kaggle)
 - Connect Network Drive for output
 - Use Checkpoint-Callbacks to store data
- Prevent logout
 - Free services often stop the python kernel after a long period of inactivity or if they are short on resources
 - If you use the free service and you need a longer runtime while you are afk you can use browser plugins to simulate activity
 - In doubt switch to a payed plan

```
from google.colab import drive
drive.mount('/content/gdrive')
```

Data Augmentation

Extend your training data



Data Augmentation

- **Definition:** Data Augmentation creates more samples by slight modifications that usually don't have an effect on the predicted class
- Example Modifications for Images
 - Shifting, Scaling, Rotation
 - Adjusting Contrast, Saturation, Brightness
 - Adding noise
- Example Modification for Audio
 - Add Pitch
 - ...
- Further Information:
 - <https://www.datacamp.com/tutorial/complete-guide-data-augmentation>

Data Augmentation with Tensorflow

- Option 1: Using special Layers in the model
 - RandomFlip
 - RandomRotation
 - With every epoch a randomly modified sample is being shown
 - We can create an extra Network for augmentation
 - These layers apply the random modification **just in training**
 - https://keras.io/api/layers/preprocessing_layers/image_augmentation/
- Option 2: Using the ImageDataGenerator
 - Here augmentation is not part of the model itself, instead it's a sperate step
 - https://www.tensorflow.org/api_docs/python/tf/keras/preprocessing/image/ImageDataGenerator
[or](#)
- Option 3: Use Keras-Function to read from Directory
 - Augmentation can be done while “streaming” data from the filesystem
 - `tf.keras.preprocessing.image_dataset_from_directory`
 - <https://keras.io/api/preprocessing/image/>

Data Augmentation Notebook

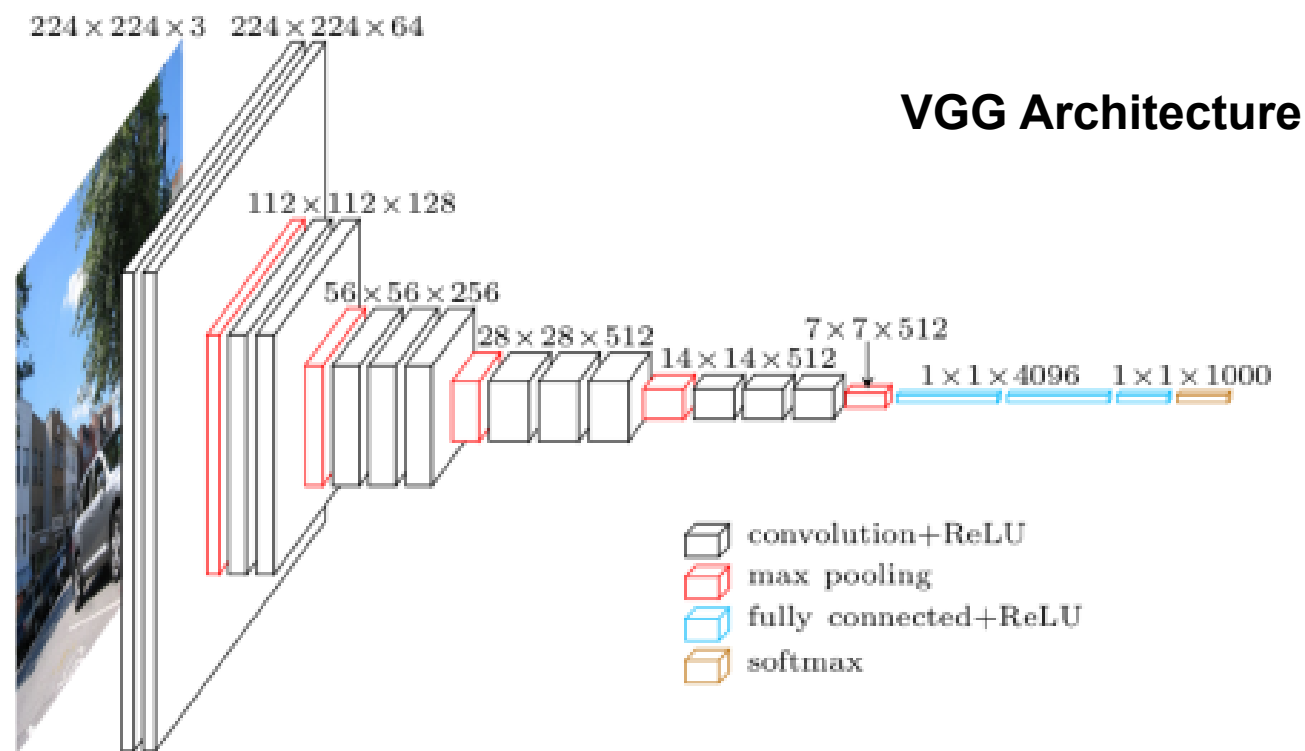
[DL_014_DataAugmentation.ipynb](#)



What did we learn from the last Notebook

- Data Augmentation helps to fight overfitting
 - **accuracy: 0.9400** - val_accuracy: **0.2145** (without augmentation)
 - **accuracy: 0.3980** - val_accuracy: **0.3218** (with augmentation)
- Networks can be nested
 - Use network as a layer to compose more complex models
 - Take care that input and output layers fit together and that input_size for first layer is given
- Streaming images with Iterators
 - Images for training don't have to be in memory, but can be loaded step by step
 - Creation of Iterators for streaming
 - `ImageDataGenerator.flow(..`
 - `tf.keras.utils.image_dataset_from_directory()`
- We have written a own training-Loop for several epochs
 - Most of the times this is not necessary but it gives you additional flexibility

CNN-Architectures and pretrained models




Pretrained Models for Classification in Keras

- Architectures
 - VGG16 - 16 Layers
 - VGG19 – 19 Layers
 - InceptionV9
 - ResNet (50 to 100Layers)
Residual Connections
 - EfficientNetB0
- Pretrained Models
 - Weights available for different image training sets
 - Imagenet, COCO etc

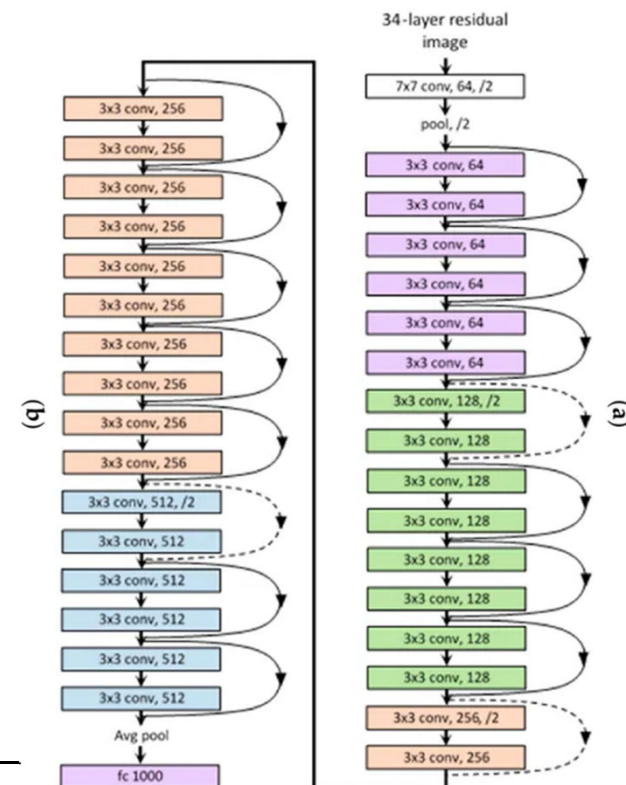
<https://pub.towardsai.net/from-vggnet-to-efficientnet-key-milestones-in-the-evolution-of-cnn-design-d778aa1e1bed>

ResNet Architecture (2015)

- ResNet is the abbreviation of Residual Neural Network
- Authors: Kaiming He et al (formerly Microsoft now Facebook)
- Design:
 - Introduction of so called “Skip-Connections” aka “residual connections” to form Residual Blocks
 - A couple of these blocks are stacked
- Effect:
 - Identity function allows the upstream gradient to better flow backwards (no vanishing gradients)
 - Networks can be deeper




Picture from Wikipedia



Picture from original Paper of Kaiming He

Applying pretrained models (Trained on ImageNet)

DL_014_VGG16.ipynb
DL_015_VGG19.ipynb
DL_016_InceptionV3.ipynb
DL_017_EfficientNetB0.ipynb



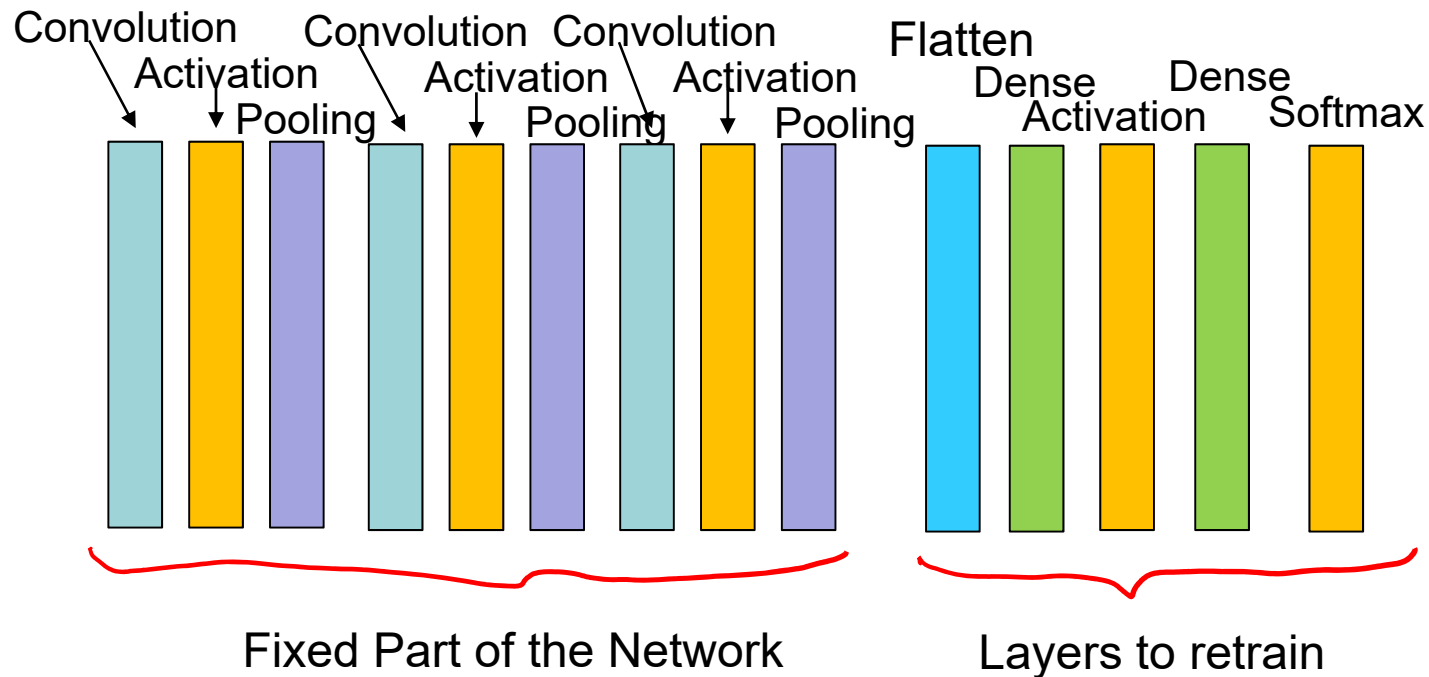
For these CNN-Architectures
we can get weights for
different image datasets

All classes of the image dataset used for training can be predicted!

Transfer learning

- NN models can just predict what they are trained for
 - Same kind of input format (pixel resolution and channels)
 - Same prediction classes (e.g 100 different classes)
- For example a NN trained on Imagenet
 - can distinguish between cats and dogs
 - cannot distinguish between different breeds of dogs
- Luckily we can reuse models trained on different tasks and extend them to new use cases
 - Early layers are frozen and weights later layers are recreated (reinitialized)
 - Then they are trained with new samples
- Transfer learning is applicable even when...
 - ...we don't have enough data for the new task to train
 - ...we don't have the time and resource to train from scratch

Transfer Learning Idea




Model is already capable of:

Detecting	Textures	Object
Edges and	& Patterns	Parts
Corners		

Transfer Learning Idea

- Different approaches possible
 - Just retrain the last Fully Connected Layer
 - Retrain multiple layers from the end (all Dense layers)
 - You can also unlock earlier Convolutional Layers and adjust weights here (e.g Fine Tune / improve results with new data)
- General Rules
 - The more layers you retain the more data and epochs are required
 - You typically shouldn't unlock certain types of layers like BatchNormalization Layers
(Reference: <https://pub.towardsai.net/batchnorm-for-transfer-learning-df17d2897db6>)

Reusing Pretrained Modells (on Imagenet) for own tasks



DL_024_Transfer Learning-EfficientNet.ipynb

<https://bit.ly/42P1jGm>